

ABTesting

October 27, 2021

1 Using A/B testing to determine whether a company should keep its old homepage

Finding out whether a company should keep its old page or its new page, or take further time to collect data and make a decision.

This project used A/B testing to determine one of the above options.

1.1 Data Wrangling

The data was collected by the company and provided by Udacity.

1.1.1 Information about the data

Converted: whether or not the user bought the product. 1 means yes, 0 means no. Landing_page: Whether the user returned to the old page or the new page. Group: Whether the user belonged to the control or the treatment group (treatment group tried out the new page). TimeStamp: Date and time. User_id: Unique ID of the user.

The data was read in below:

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
```

```
[2]: ab_data = pd.read_csv('ab_data.csv')
ab_data.head()
```

```
[2]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

Obtaining basic information about the dataset. This information is listed in the next section:

```
[3]: ab_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id          294478 non-null  int64
1   timestamp        294478 non-null  object
2   group            294478 non-null  object
3   landing_page     294478 non-null  object
4   converted        294478 non-null  int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB

```

```
[4]: ab_data.loc[(ab_data['group']=='treatment') & (ab_data['landing_page'] == 'old_page')]
```

```

[4]:      user_id      timestamp      group landing_page  converted
308    857184  2017-01-20 07:34:59.832626  treatment    old_page         0
327    686623  2017-01-09 14:26:40.734775  treatment    old_page         0
357    856078  2017-01-12 12:29:30.354835  treatment    old_page         0
685    666385  2017-01-23 08:11:54.823806  treatment    old_page         0
713    748761  2017-01-10 15:47:44.445196  treatment    old_page         0
...
293773  688144  2017-01-16 20:34:50.450528  treatment    old_page         1
293817  876037  2017-01-17 16:15:08.957152  treatment    old_page         1
293917  738357  2017-01-05 15:37:55.729133  treatment    old_page         0
294014  813406  2017-01-09 06:25:33.223301  treatment    old_page         0
294252  892498  2017-01-22 01:11:10.463211  treatment    old_page         0

```

[1965 rows x 5 columns]

So 1965 rows are incorrectly labelled so that “treatment” is mapped to “old_page”.

Checking whether control users were given “old_page”:

```
[5]: ab_data.loc[(ab_data['group']=='control') & (ab_data['landing_page'] == 'new_page')]
```

```

[5]:      user_id      timestamp      group landing_page  converted
22    767017  2017-01-12 22:58:14.991443  control    new_page         0
240   733976  2017-01-11 15:11:16.407599  control    new_page         0
490   808613  2017-01-10 21:44:01.292755  control    new_page         0
846   637639  2017-01-11 23:09:52.682329  control    new_page         1
850   793580  2017-01-08 03:25:33.723712  control    new_page         1
...
293894  741581  2017-01-09 20:49:03.391764  control    new_page         0
293996  942612  2017-01-08 13:52:28.182648  control    new_page         0
294200  928506  2017-01-13 21:32:10.491309  control    new_page         0
294253  886135  2017-01-06 12:49:20.509403  control    new_page         0

```

```
294331    689637    2017-01-13 11:34:28.339532    control    new_page    0
```

```
[1928 rows x 5 columns]
```

So 1928 rows are incorrectly labelled.

Finding how many unique users there were:

```
[6]: ab_data.nunique()
```

```
[6]: user_id      290584
      timestamp    294478
      group        2
      landing_page  2
      converted     2
      dtype: int64
```

```
[7]: ab_data.loc[(ab_data['converted'] == 1)]
```

```
[7]:
```

	user_id	timestamp	group	landing_page	converted
4	864975	2017-01-21 01:52:26.210827	control	old_page	1
6	679687	2017-01-19 03:26:46.940749	treatment	new_page	1
8	817355	2017-01-04 17:58:08.979471	treatment	new_page	1
9	839785	2017-01-15 18:11:06.610965	treatment	new_page	1
15	644214	2017-01-22 02:05:21.719434	control	old_page	1
...
294396	838593	2017-01-15 09:56:31.455023	treatment	new_page	1
294405	712217	2017-01-11 10:34:30.176801	control	old_page	1
294420	795742	2017-01-09 01:06:58.299207	control	old_page	1
294430	733871	2017-01-21 17:54:08.810964	treatment	new_page	1
294443	665217	2017-01-10 23:29:01.767720	control	old_page	1

```
[35237 rows x 5 columns]
```

```
[8]: ab_data.isnull().sum()
```

```
[8]: user_id      0
      timestamp    0
      group        0
      landing_page  0
      converted     0
      dtype: int64
```

Drop the incorrectly labelled data:

```
[9]: treatment = ab_data[(ab_data['group']=='treatment') & (ab_data['landing_page']_
      ↪ == 'new_page')]
```

```
[10]: control = ab_data[(ab_data['group']=='control') & (ab_data['landing_page'] ==_
    ↪ 'old_page')]
```

```
[11]: ab_cleaned = pd.concat([treatment,control], ignore_index='True')
```

```
[12]: ab_cleaned.head()
```

```
[12]:
```

	user_id	timestamp	group	landing_page	converted
0	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
1	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
2	679687	2017-01-19 03:26:46.940749	treatment	new_page	1
3	817355	2017-01-04 17:58:08.979471	treatment	new_page	1
4	839785	2017-01-15 18:11:06.610965	treatment	new_page	1

```
[13]: ab_cleaned.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 290585 entries, 0 to 290584
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id          290585 non-null  int64
1   timestamp        290585 non-null  object
2   group            290585 non-null  object
3   landing_page     290585 non-null  object
4   converted        290585 non-null  int64
dtypes: int64(2), object(3)
memory usage: 11.1+ MB
```

```
[14]: ab_cleaned.nunique()
```

```
[14]: user_id          290584
timestamp          290585
group               2
landing_page       2
converted          2
dtype: int64
```

The number of unique users is 290584.

```
[15]: print(ab_cleaned[ab_cleaned.duplicated()])
```

```
Empty DataFrame
Columns: [user_id, timestamp, group, landing_page, converted]
Index: []
```

```
[16]: ab_cleaned.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 290585 entries, 0 to 290584
```

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	user_id	290585 non-null	int64
1	timestamp	290585 non-null	object
2	group	290585 non-null	object
3	landing_page	290585 non-null	object
4	converted	290585 non-null	int64

dtypes: int64(2), object(3)

memory usage: 11.1+ MB

```
[17]: N_conversions = ab_cleaned[(ab_cleaned['converted'] == 1)].count()
      N_rows = ab_cleaned.count()
```

```
[18]: p_convert = N_conversions/N_rows
      p_convert
```

```
[18]: user_id      0.119597
      timestamp    0.119597
      group        0.119597
      landing_page  0.119597
      converted     0.119597
      dtype: float64
```

Now calculating the probability of converting given that the individual was in the control group:

```
[19]: control = ab_cleaned[ab_cleaned['group']=='control'].count()
      p_old = ab_cleaned[ab_cleaned['group']=='control']['converted'].sum()/control
      p_old
```

```
[19]: user_id      0.120386
      timestamp    0.120386
      group        0.120386
      landing_page  0.120386
      converted     0.120386
      dtype: float64
```

```
[20]: control
```

```
[20]: user_id      145274
      timestamp    145274
      group        145274
      landing_page  145274
      converted     145274
      dtype: int64
```

```
[21]: treatment = ab_cleaned[ab_cleaned['group']=='treatment'].count()
```

```
p_new = ab_cleaned[ab_cleaned['group']=='treatment']['converted'].sum()/
↳treatment
p_new
```

```
[21]: user_id      0.118807
      timestamp   0.118807
      group       0.118807
      landing_page 0.118807
      converted    0.118807
      dtype: float64
```

```
[22]: new_page = ab_cleaned[ab_cleaned['landing_page']=='new_page'].count()
      new_page
```

```
[22]: user_id      145311
      timestamp   145311
      group       145311
      landing_page 145311
      converted    145311
      dtype: int64
```

```
[23]: p_new_page = new_page['landing_page']/N_rows
      p_new_page
```

```
[23]: user_id      0.500064
      timestamp   0.500064
      group       0.500064
      landing_page 0.500064
      converted    0.500064
      dtype: float64
```

```
[24]: obs_diff = abs(p_new - p_old)
      obs_diff
```

```
[24]: user_id      0.001579
      timestamp   0.001579
      group       0.001579
      landing_page 0.001579
      converted    0.001579
      dtype: float64
```

1.1.2 Does the new treatment group have a higher probability of converting?

The previous calculations show that the dataset is split 50/50 into treatment and control data.

The probability of converting on the old page is 0.1549% greater.

The probability difference requires further investigation as it is a difference, but it is not clear whether we can be confident it is statistically significant.

1.1.3 Hypothesis testing

Hypothesis O: The pages have the same conversion rates, so that $p_{\text{new}} = p_{\text{old}} = p_{\text{convert}}$

Hypothesis 1: The old page has a significantly higher conversion rate, with a confidence interval of 95%.

Now we simulate 10000 samples and calculate the probability differences for each one:

```
[25]: treatment['landing_page']
```

```
[25]: 145311
```

```
[26]: p_new = p_new['landing_page']
```

```
[27]: N_new = treatment['landing_page']  
new_converted_simulation = np.random.binomial(N_new,p_new,10000)/N_new  
new_converted_simulation
```

```
[27]: array([0.11819477, 0.11864897, 0.11817412, ..., 0.11782315, 0.12006662,  
          0.1187866 ])
```


```
[28]: N_old = control['landing_page']  
p_old = p_old['landing_page']  
old_converted_simulation = np.random.binomial(N_old,p_old,10000)/N_old  
old_converted_simulation
```

```
[28]: array([0.12141884, 0.1208475 , 0.11837631, ..., 0.12051021, 0.11830747,  
          0.1209301 ])
```

```
[29]: p_diffs = old_converted_simulation - new_converted_simulation  
obs_diff = obs_diff['landing_page']
```

Plotting a histogram of the results:

```
[30]: plt.hist(p_diffs)  
plt.axvline(x=obs_diff,color='r')  
plt.ylabel('Frequency Density')  
plt.xlabel('Probability Difference')  
plt.show()
```



The above graph shows that it is very likely the null hypothesis should be rejected, as about half the values simulated are above the observed difference.

Now calculating what proportion of the simulated probability differences are greater than the observed probability difference calculated earlier. This will confirm or deny H_0 .

```
[31]: obs_diff
```

```
[31]: 0.0015790565976871451
```

```
[57]: proportion = (p_diffs > obs_diff).sum()
      proportion_value = proportion/len(p_diffs)
```

```
[58]: proportion_value
```

```
[58]: 0.4965
```

Here we see that nearly 50% of the simulated probability differences are greater than the observed difference. This suggests that the observed difference lies close to the center of the probability difference distribution and therefore is not statistically significant.

However it is also clear that the mean probability difference is greater than 0, which seems to contradict the null hypothesis.

1.1.4 Computing the Z-score using statsmodel.api

```
[34]: import statsmodels.api as sm

import numpy as np
from statsmodels.stats.proportion import proportions_ztest

convert_new = ab_cleaned[ab_cleaned['group']=='treatment']['converted'].sum()
convert_old = ab_cleaned[ab_cleaned['group']=='control']['converted'].sum()

n_old = control['landing_page']
n_new = treatment['landing_page']

# array with the numbers of converted customers in both control and treatment_
→ groups
count = np.array([convert_old, convert_new])
# total number of observations in each group
number_of_obs = np.array([n_old, n_new])
# Calculates the p-value for these two samples
stat, pval = proportions_ztest(count, number_of_obs, alternative='larger')
```

```
[35]: print(pval)
```

0.09482629485940902

The p-value here is 9.48%. This is much greater than the significance level chosen of 5%. Therefore the result here is *not* statistically significant and the null hypothesis, the conversion rates for both old and new pages are equal, is accepted.

1.2 Performing an A/B test using regression

As this is a classification problem, logistic regression was used.

Some pre-processing was necessary:

1. A new column 'intercept' was created.
2. A column 'ab_page' which essentially encodes the treatment as '1' and the control as '0'.

```
[36]: ab_cleaned['intercept'] = 1
ab_cleaned.head()
```

```
[36]:
```

	user_id	timestamp	group	landing_page	converted	\
0	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
1	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
2	679687	2017-01-19 03:26:46.940749	treatment	new_page	1	
3	817355	2017-01-04 17:58:08.979471	treatment	new_page	1	
4	839785	2017-01-15 18:11:06.610965	treatment	new_page	1	

intercept

0	1
1	1
2	1
3	1
4	1

```
[37]: from sklearn.preprocessing import LabelBinarizer
lb = LabelBinarizer()
ab_cleaned['ab_page'] = lb.fit_transform(ab_cleaned['group'])
```

```
[38]: ab_cleaned.head()
```

```
[38]:
```

	user_id	timestamp	group	landing_page	converted	\
0	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
1	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
2	679687	2017-01-19 03:26:46.940749	treatment	new_page	1	
3	817355	2017-01-04 17:58:08.979471	treatment	new_page	1	
4	839785	2017-01-15 18:11:06.610965	treatment	new_page	1	

	intercept	ab_page
0	1	1
1	1	1
2	1	1
3	1	1
4	1	1

```
[39]: ab_cleaned.tail()
```

```
[39]:
```

	user_id	timestamp	group	landing_page	converted	\
290580	718310	2017-01-21 22:44:20.378320	control	old_page	0	
290581	751197	2017-01-03 22:28:38.630509	control	old_page	0	
290582	945152	2017-01-12 00:51:57.078372	control	old_page	0	
290583	734608	2017-01-22 11:45:03.439544	control	old_page	0	
290584	697314	2017-01-15 01:20:28.957438	control	old_page	0	

	intercept	ab_page
290580	1	0
290581	1	0
290582	1	0
290583	1	0
290584	1	0

Using logistic regression model to fit the data:

```
[55]: x = ab_cleaned[['intercept', 'ab_page']]
y = ab_cleaned['converted']
model = sm.Logit(y,x)
results = model.fit()
```

```

Optimization terminated successfully.
      Current function value: 0.366118
      Iterations 6

```

```
[56]: results.summary()
```

```

[56]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                Logit Regression Results
=====
Dep. Variable:                converted    No. Observations:                290585
Model:                        Logit       Df Residuals:                  290583
Method:                       MLE        Df Model:                      1
Date:                         Wed, 27 Oct 2021    Pseudo R-squ.:                8.085e-06
Time:                         09:26:02    Log-Likelihood:               -1.0639e+05
converged:                     True        LL-Null:                     -1.0639e+05
Covariance Type:              nonrobust    LLR p-value:                  0.1897
=====
                                coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept          -1.9888         0.008   -246.669      0.000      -2.005      -1.973
ab_page            -0.0150         0.011    -1.312      0.190      -0.037       0.007
=====
      """

```

1.2.1 Conclusions based on this Logistic Regression model:

The p-value of this model is 19.0%. This is much higher than the significance level of 5%. It means that the difference in probabilities of conversion is statistically insignificant.

In this case, the p-value calculated is different because the logistic regression model uses a sigma function as the probability distribution, whereas calculating z-scores relies on the normal distribution. These are two very differently shaped distributions and they produce different results.

1.2.2 What is the impact of the user's country, if there is any impact, on the conversion rate?

A new set of data had to be imported and joined to the cleaned A/B testing data on the column 'user_id'.

This was then used to estimate the impact of the user's country on the conversion rate.

```
[59]: countries = pd.read_csv('countries.csv')
```

```
[60]: countries.head()
```

```

[60]:   user_id  country
0    834778      UK
1    928468      US
2    822059      UK

```

```
3    711597    UK
4    710616    UK
```

```
[61]: df_merged = pd.merge(countries,ab_cleaned, how='left', on=['user_id'])
```

```
[62]: df_merged.head()
```

```
[62]:
```

	user_id	country	timestamp	group	landing_page \
0	834778	UK	2017-01-14 23:08:43.304998	control	old_page
1	928468	US	2017-01-23 14:44:16.387854	treatment	new_page
2	822059	UK	2017-01-16 14:04:14.719771	treatment	new_page
3	711597	UK	2017-01-22 03:14:24.763511	control	old_page
4	710616	UK	2017-01-16 13:14:44.000513	treatment	new_page

	converted	intercept	ab_page
0	0	1	0
1	0	1	1
2	1	1	1
3	0	1	0
4	0	1	1

```
[67]: df_merged.tail()
```

```
[67]:
```

	user_id	country	timestamp	group	landing_page \
290580	653118	0	2017-01-09 03:12:31.034796	control	old_page
290581	878226	0	2017-01-05 15:02:50.334962	control	old_page
290582	799368	0	2017-01-09 18:07:34.253935	control	old_page
290583	655535	1	2017-01-09 13:30:47.524512	treatment	new_page
290584	934996	0	2017-01-09 00:30:08.377677	control	old_page

	converted	intercept	ab_page
290580	0	1	0
290581	0	1	0
290582	0	1	0
290583	0	1	1
290584	0	1	0

Turning the categorical data for 'countries' into binary values:

```
[64]: df_merged['country'] = pd.get_dummies(df_merged['country'])
```

```
[66]: df_merged['country']
```

```
[66]:
```

	country
0	0
1	0
2	0
3	0
4	0

```

..
290580    0
290581    0
290582    0
290583    1
290584    0
Name: country, Length: 290585, dtype: uint8

```

```
[65]: df_merged.head()
```

```

[65]:   user_id  country          timestamp    group landing_page \
0    834778        0  2017-01-14 23:08:43.304998  control    old_page
1    928468        0  2017-01-23 14:44:16.387854  treatment  new_page
2    822059        0  2017-01-16 14:04:14.719771  treatment  new_page
3    711597        0  2017-01-22 03:14:24.763511  control    old_page
4    710616        0  2017-01-16 13:14:44.000513  treatment  new_page

   converted  intercept  ab_page
0          0          1         0
1          0          1         1
2          1          1         1
3          0          1         0
4          0          1         1

```

Now fitting the logistic regression model, but with two independent variables instead of one:

```

[68]: model2 = sm.Logit(df_merged['converted'],df_merged[['intercept',
↪'ab_page','country']])
results = model2.fit()

```

```

Optimization terminated successfully.
      Current function value: 0.366113
      Iterations 6

```

```
[69]: results.summary()
```

```

[69]: <class 'statsmodels.iolib.summary.Summary'>
      """

                                Logit Regression Results
=====
Dep. Variable:                  converted    No. Observations:                  290585
Model:                            Logit      Df Residuals:                  290582
Method:                            MLE        Df Model:                      2
Date:                Wed, 27 Oct 2021    Pseudo R-squ.:                  2.065e-05
Time:                   09:41:33      Log-Likelihood:                 -1.0639e+05
converged:                            True      LL-Null:                   -1.0639e+05
Covariance Type:                nonrobust    LLR p-value:                   0.1112
=====

```

	coef	std err	z	P> z	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
intercept	-1.9867	0.008	-243.359	0.000	-2.003	-1.971
ab_page	-0.0150	0.011	-1.309	0.191	-0.037	0.007
country	-0.0434	0.027	-1.627	0.104	-0.096	0.009
=====	=====	=====	=====	=====	=====	=====
"""						

1.2.3 Conclusions:

The model provides a p-value of 11.1%.

The null hypothesis is that country has no effect on the result. This was tested with a significance level of 5%. It is a right-tailed test.

The outcome is that the calculated p-value is greater than the significance level of 5%.

This leads us to conclude that country does not have a statistically significant effect on the conversion rate.

1.2.4 Other possible factors which should be accounted for:

There are some limitations of this model.

Its result differs partly because it assumes a logistic regression rather than a normal distribution. The functions employed in both cases are quite different.

Further questions to be explored:

- Do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time?
- How long do you run to render a decision that neither page is better than another?